

Time Series Analysis and Recurrent Neural Network

Giacomo Barzon - 3626438

Exercise 4

November 20, 2019

1 Task 1

1.0.1 1) Create a 2-variate AR(2) time series

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import chi2
from scipy.stats import f

In [2]: # INPUT PARAMETERS
# total time steps
T = 1000

# set random seed
np.random.seed(19062019)

# define empty ts matrix
data = np.zeros((2,T))

# setting initial conditions
data[:,0] = [0., 0.]

# model parameters (a0 = [0,0])
A1 = np.array(( [0.2, -0.2], [0., 0.1] ))
A2 = np.array(( [0.1, -0.1], [0., 0.1] ))

A_model = np.concatenate( ([0.0,0.0], A1, A2), axis=None)

# noise covariance matrix
cov = np.array(( [0.01, 0], [0, 0.01] ))

In [3]: # CREATE THE TIME SERIES
# generate values from the gaussian noise
epsilon = np.random.multivariate_normal([0.,0.], cov, T).T
```

```

# loop over time steps
data[:, 1] = np.dot(A1, data[:,0]) + epsilon[:,0]

for i in range(2,T):
    data[:, i] = np.dot(A1, data[:,i-1]) + np.dot(A2, data[:,i-2]) + epsilon[:,i-1]

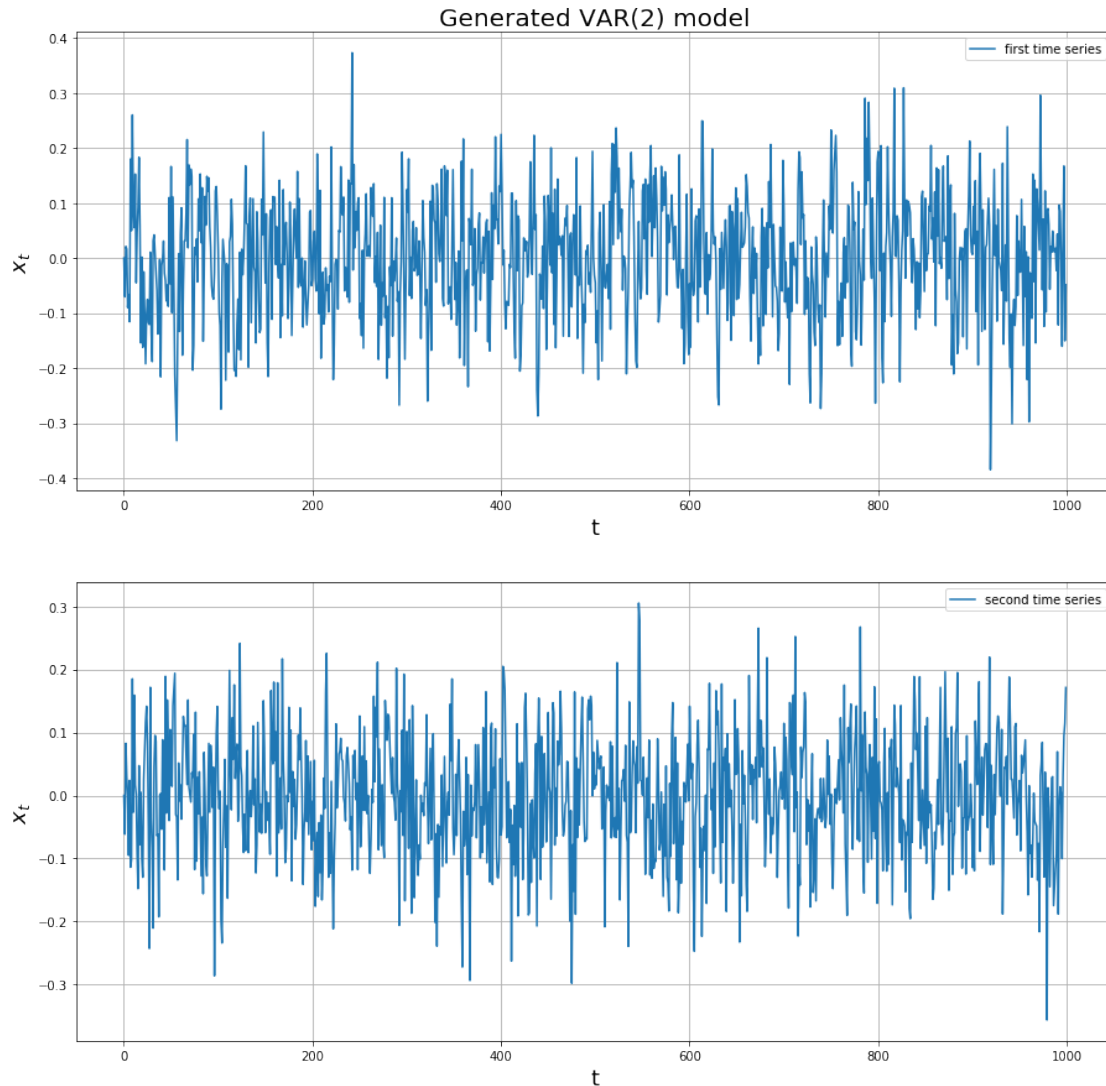
In [4]: # PLOT GENERATED TIME SERIES
fig1 = plt.subplots(figsize=[15,15])

plt.subplot(2,1,1)
plt.plot(np.arange(T), data[0,:], label='first time series')
plt.xlabel('t', fontsize = 18)
plt.ylabel(r'$x_t$', fontsize = 18)
plt.grid()
plt.legend()
plt.title('Generated VAR(2) model', fontsize = 20)

plt.subplot(2,1,2)
plt.plot(np.arange(T), data[1,:], label='second time series')
plt.xlabel('t', fontsize = 18)
plt.ylabel(r'$x_t$', fontsize = 18)
plt.grid()
plt.legend()

plt.show(fig1)

```



1.0.2 2) Given your knowledge of model parameters above, does x_1 Granger-cause x_2 , or does x_2 Granger-cause x_1 according to Granger's definition? Why?

As we can notice by looking at the matrix of the generating model, the lower-left (2,1) coefficients are both null: this means that the generated values of x_2 don't depend on the previous value of x_1 and for this reason we can deduce that x_1 does not Granger-cause x_2 . On the other hand the upper-left (1,2) are different from zero, so we can deduce that x_2 Granger-cause x_1 .

1.0.3 3) Assuming that the order of the model ($p = 2$) is known, use the log-likelihood-ratio test statistic to confirm your results from (2)

```
In [5]: # USEFUL FUNCTIONS
        # Coefficient estimation for AR(p) model
        def coeff_AR(data, p):
```

```

xT = data[p:]
Xp = np.zeros((len(data)-p,p+1))

# add 1 as first element
Xp[:,0] = 1

# fill matrix Xp with data
for i in range(1,p+1):
    Xp[:,i] = data[p-i:-i]

# compute coefficients
temp = np.matmul(np.transpose(Xp),Xp)
temp = np.linalg.inv(temp)
a = np.matmul(np.transpose(Xp),xT)
a = np.matmul(temp,a)

return a

# define p as a global variable
p = 2

# coefficient estimation for VAR(p)
def coeff_VAR(data, p):
    data = data.T
    T = data.shape[0]
    n = data.shape[1]

    # define empty matrix
    A = np.zeros(( n, n*p+1 ))
    xT = np.zeros(T-p)
    Xp = np.zeros(( T-p), n*p+1 ))

    # add column of 1s
    Xp[:,0] = np.zeros(T-p) + 1

    # fill matrix Xp with data
    for i in range(n):
        for k in range(1,p+1):
            Xp[ :, i*p +k ] = data[p-k:-k, i]

    # compute coefficients
    for i in range(n):
        xT = data[p:, i]

        temp = np.matmul(np.transpose(Xp),Xp)
        temp = np.linalg.inv(temp)
        a = np.matmul(np.transpose(Xp),xT)
        a = np.matmul(temp,a)

```

```

        A[i] = a

    return A

# inferred VAR time series, given coefficient matrices
def pred_VAR(data, a0, A1, A2):
    T = data.shape[1]
    n = data.shape[0]

    # define empty array
    x_new = np.zeros((n,T))
    # set first p values
    x_new[:, :p] = data[:, :p]

    for i in range(p,T):
        x_new[:, i] = a0 + np.dot(A1, data[:,i-1]) + np.dot(A2, data[:,i-2])

    return x_new

# Residuals of the VAR(p) model
def res(data, a0, A1, A2):
    return pred_VAR(data, a0, A1, A2)[:,p:] - data[:,p:]

# Covariance matrix of the residuals
def sigma(data, a0, A1, A2):
    return np.cov(res(data, a0, A1, A2))

# Log-likelihood for VAR(p) model
def log_likelihood(data, a0, A1, A2):
    T = data.shape[1]
    n = data.shape[0]

    res1 = res(data, a0, A1, A2)
    sigma1 = sigma(data, a0, A1, A2)
    det = np.linalg.det(sigma1)

    t1 = - 0.5 * (T-p)*n * np.log(2*np.pi)
    t2 = - 0.5 * (T-p) * np.log(det)
    t3 = 0.
    for i in range(T-p):
        temp = np.dot(np.linalg.inv(sigma1), res1[:,i])
        t3 += np.dot(res1[:,i], temp)
    t3 = - 0.5 * t3

    return t1 + t2 + t3

```

```

In [6]: # NULL HYPOTHESIS: no Granger Causalities
        # diagonal A1 and A2 matrices / two sigle AR model

```

```

a1 = coeff_AR(data[0,:], p=2)
a2 = coeff_AR(data[1,:], p=2)

A0_null = np.array([a1[0], a2[0]])
A1_null = np.array([a1[1], 0.0], [0.0, a2[1]])
A2_null = np.array([a1[2], 0.0], [0.0, a2[2]])

A_null = np.concatenate( (A0_null, A1_null, A2_null), axis=None)

print('Null hypothesis parameters')
print('A0:\n', A0_null)
print('A1:\n', A1_null)
print('A2:\n', A2_null)

```

```

Null hypothesis parameters
A0:
 [ 9.90742761e-05 -1.68220466e-03]
A1:
 [[0.24370888 0.          ]
 [0.          0.13194647]]
A2:
 [[0.09095963 0.          ]
 [0.          0.10616485]]

```

```

In [7]: # COMPLETE HYPOTHESIS: VAR(2) model
A = coeff_VAR(data, p=2)

A0_var = A[:,0]
A1_var = np.array((A[:,1], A[:,3])).T
A2_var = np.array((A[:,2], A[:,4])).T

A_var = np.concatenate( (A0_var, A1_var, A2_var), axis=None)

print('Complete model parameters')
print('A0:\n', A0_var)
print('A1:\n', A1_var)
print('A2:\n', A2_var)

```

```

Complete model parameters
A0:
 [-0.00066633 -0.00167826]
A1:
 [[ 0.22035789 -0.17975054]
 [ 0.00785193  0.13189426]]
A2:
 [[ 0.09273513 -0.11843226]
 [-0.00411085  0.10766649]]

```

```
In [8]: # FIRST ALTERNATIVE HYPOTHESIS: x1 doesn't Granger causes x2
```

```
A0_x2 = np.array((A[0,0], a2[0]))
A1_x2 = np.array((A1_var))
A1_x2[1] = np.array((0.0, a2[1]))
A2_x2 = np.array((A2_var))
A2_x2[1] = np.array((0.0, a2[2]))

A_x2 = np.concatenate( (A0_x2,A1_x2,A2_x2), axis=None)

print('(x1 not G.C. x2) model parameters')
print('A0:\n', A0_x2)
print('A1:\n', A1_x2)
print('A2:\n', A2_x2)
```

```
(x1 not G.C. x2) model parameters
```

```
A0:
```

```
[-0.00066633 -0.0016822 ]
```

```
A1:
```

```
[[ 0.22035789 -0.17975054]
 [ 0.          0.13194647]]
```

```
A2:
```

```
[[ 0.09273513 -0.11843226]
 [ 0.          0.10616485]]
```

```
In [9]: # SECOND ALTERNATIVE HYPOTHESIS: x2 doesn't Granger causes x1
```

```
A0_x1 = np.array((A[0,0], a2[0]))
A1_x1 = np.array((A1_var))
A1_x1[0] = np.array((a1[1], 0.0))
A2_x1 = np.array((A2_var))
A2_x1[0] = np.array((a1[2], 0.0))

A_x1 = np.concatenate( (A0_x1,A1_x1,A2_x1), axis=None)

print('(x2 not G.C. x1) model parameters')
print('A0:\n', A0_x1)
print('A1:\n', A1_x1)
print('A2:\n', A2_x1)
```

```
(x2 not G.C. x1) model parameters
```

```
A0:
```

```
[-0.00066633 -0.0016822 ]
```

```
A1:
```

```
[[0.24370888 0.          ]
 [0.00785193 0.13189426]]
```

```
A2:
```

```
[[ 0.09095963 0.          ]
 [-0.00411085 0.10766649]]
```

```

In [10]: # PLOTS OF THE COEFFICIENTS FOR DIFFERENT MODELS
A = np.array((A_var, A_null, A_x2, A_x1))

fig1 = plt.subplots(figsize=[15,15])
width = 0.35
xs = np.arange(A.shape[1])

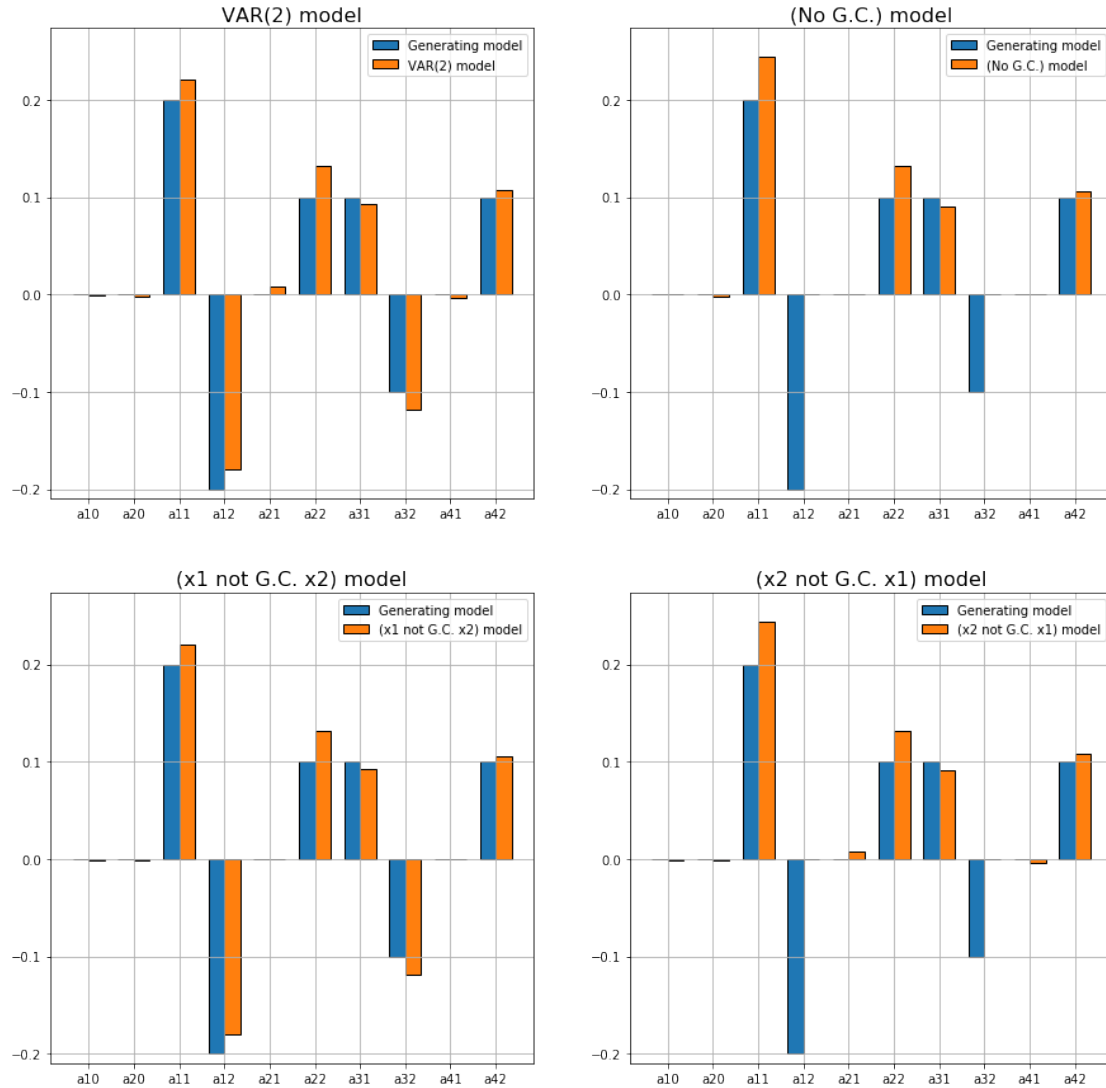
x_ticks = ['a10', 'a20', 'a11', 'a12', 'a21', 'a22', 'a31', 'a32', 'a41', 'a42']
titles = ['VAR(2) model', '(No G.C.) model', '(x1 not G.C. x2) model', '(x2 not G.C. x1) model']

for i in range(len(A)):
    plt.subplot(2,2,i+1)

    plt.bar(xs - width/2, A_model[:,i], width, label='Generating model', edgecolor = 'black')
    plt.bar(xs + width/2, A[i,:], width, label=titles[i], edgecolor = 'black')
    plt.xticks( xs, x_ticks )
    plt.ylim( [A.min()-0.03, A.max()+0.03] )
    plt.grid()
    plt.legend()
    plt.title(titles[i], fontsize = 16)

plt.show(fig1)

```

```
In [11]: # LOG-LIKELIHOOD COMPUTATION FOR DIFFERENT MODELS
```

```
k = 2
```

```
ll = log_likelihood(data, A0_var, A1_var, A2_var)
```

```
print('log-likelihood for complete model:', ll)
```

```
llnull = log_likelihood(data, A0_null, A1_null, A2_null)
```

```
print('log-likelihood for restricted model:', llnull)
```

```
llx1 = log_likelihood(data, A0_x1, A1_x1, A2_x1)
```

```
print('log-likelihood for (x2 not G.C. x1) model:', llx1)
```

```
llx2 = log_likelihood(data, A0_x2, A1_x2, A2_x2)
```

```
print('log-likelihood for (x1 not G.C. x2) model:', llx2)
```

```

log-likelihood for complete model:, 1786.56817554403
log-likelihood for restricted model: 1764.08138802638
log-likelihood for (x2 not G.C. x1) model: 1764.096829255153
log-likelihood for (x1 not G.C. x2) model: 1786.528964484111

```

```
In [12]: # LLR-TEST - respect to the Complete VAR model
```

```
from scipy.stats import f
```

```
D1 = -2. * (llx1 - ll)
```

```
D2 = -2. * (llx2 - ll)
```

```
print('D for (x2 not G.C. x1) model:', D1)
```

```
print('cdf(D) for (x2 not G.C. x1) model:', chi2.cdf(D1, k) )
```

```
print()
```

```
print('D for (x1 not G.C. x2) model:', D2)
```

```
print('cdf(D) for (x1 not G.C. x2) model:', chi2.cdf(D2, k) )
```

```
D for (x2 not G.C. x1) model: 44.942692577754315
```

```
cdf(D) for (x2 not G.C. x1) model: 0.9999999998258922
```

```
D for (x1 not G.C. x2) model: 0.07842211983825109
```

```
cdf(D) for (x1 not G.C. x2) model: 0.03845225646081886
```

```
In [13]: # F-TEST - respect to the Complete VAR model
```

```
from scipy.stats import f
```

```
sigma_var = sigma(data, A0_var, A1_var, A2_var)
```

```
sigma_x1 = sigma(data,A0_x1, A1_x1, A2_x1)
```

```
sigma_x2 = sigma(data,A0_x2, A1_x2, A2_x2)
```

```
Fx1 = (T-p-1) / k * ( np.log(np.linalg.det(sigma_x1)) - np.log(np.linalg.det(sigma_var))
```

```
Fx2 = (T-p-1) / k * ( np.log(np.linalg.det(sigma_x2)) - np.log(np.linalg.det(sigma_var))
```

```
print('F for (x2 not G.C. x1) model:', Fx1)
```

```
print('cdf(F) for (x2 not G.C. x1) model:', f.cdf(Fx1, k, (T-p-1)-k*(p+1)-1 ) )
```

```
print()
```

```
print('F for (x1 not G.C. x2) model:', Fx2)
```

```
print('cdf(F) for (x1 not G.C. x2) model:', f.cdf(Fx2, k, (T-p-1)-k*(p+1)-1 ) )
```

```
F for (x2 not G.C. x1) model: 22.42235108827798
```

```
cdf(F) for (x2 not G.C. x1) model: 0.9999999997006309
```

```
F for (x1 not G.C. x2) model: 0.03917177028048702
```

```
cdf(F) for (x1 not G.C. x2) model: 0.03841298654551087
```

```
In [14]: # LLR-TEST - respect to the Complete VAR model
```

```
D1 = -2. * (llnull - llx1)
```

```

D2 = -2. * (llnull - llx2)

print('D for (x2 not G.C. x1) model:', D1)
print('cdf(D) for (x2 not G.C. x1) model:', chi2.cdf(D1, k) )
print()
print('D for (x1 not G.C. x2) model:', D2)
print('cdf(D) for (x1 not G.C. x2) model: ', chi2.cdf(D2, k) )

D for (x2 not G.C. x1) model: 0.0308824575458857
cdf(D) for (x2 not G.C. x1) model: 0.015322624251171488

D for (x1 not G.C. x2) model: 44.89515291546195
cdf(D) for (x1 not G.C. x2) model: 0.9999999998217041

In [15]: # F-TEST - respect to the 2-AR model
sigma_null = sigma(data, A0_null, A1_null, A2_null)

Fx1 = (T-p-1) / k * ( np.log(np.linalg.det(sigma_null)) - np.log(np.linalg.det(sigma_x1)) )
Fx2 = (T-p-1) / k * ( np.log(np.linalg.det(sigma_null)) - np.log(np.linalg.det(sigma_x2)) )

print('F for (x2 not G.C. x1) model:', Fx1)
print('cdf(F) for (x2 not G.C. x1) model:', f.cdf(Fx1, k, (T-p-1)-k*(p+1)-1 ) )
print()
print('F for (x1 not G.C. x2) model:', Fx2)
print('cdf(F) for (x1 not G.C. x2) model:', f.cdf(Fx2, k, (T-p-1)-k*(p+1)-1 ) )

F for (x2 not G.C. x1) model: 0.04190457815246251
cdf(F) for (x2 not G.C. x1) model: 0.0410370170910969

F for (x1 not G.C. x2) model: 22.425083896149957
cdf(F) for (x1 not G.C. x2) model: 0.9999999997014126

```

By looking at the value of the F test respect to the estimated VAR(2) model, we can reject the assumption that x2 Granger-cause x1 at 1% significance level, while we can accept the assumption that x1 Granger-cause x2 at 1% significance level.

2 Task 2

2.0.1 Consider a linear Gaussian state space model:

$$z_t = Az_{t-1} + \epsilon, \epsilon \sim N(0, \Sigma)$$

$$x_t = Bz_t + \eta, \eta \sim N(0, \Gamma)$$

2.0.2 Derive the M-step for the latent state noise Σ and the observation matrix B by maximizing the expected log-likelihood $E[\log p(X, Z)]$, with respect to Σ and B , where $X = \{x_t \mid t \in 1 \dots T\}$ and $Z = \{z_t \mid t \in 1 \dots T\}$ are the sets of all latent states and observations from time 1 to T

$$ELBO[q, \theta] = E_q[\log p_\theta(x, z)] + H[q(z)]$$

Since in the M-step we maximize respect to the different parameters θ , the proposal density distribution q is considered fixed, so we take into account only the first term from above. From here on out we ignore that the expectation value is respect to the proposal density distribution $q(z)$ and that the probability depends on the different parameter θ .

Then, we can expand the first term as:

$$\begin{aligned} E[\log p(x, z)] &= E[\log p(z_1)] + E\left[\sum_{t=2}^T \log p(z_t | z_{t-1})\right] + E\left[\sum_{t=1}^T \log p(x_t | z_t)\right] \\ &= (1) + (2) + (3) \end{aligned}$$

where:

$$\begin{aligned} (1) &= E\left[-\frac{1}{2} \log |\Sigma| - \frac{1}{2} (z_1 - \mu_0)^T \Sigma^{-1} (z_1 - \mu_0)\right] \\ (2) &= E\left[-\frac{T-1}{2} \log |\Sigma| - \frac{1}{2} \sum_{t=2}^T (z_t - A z_{t-1})^T \Sigma^{-1} (z_t - A z_{t-1})\right] \\ (3) &= E\left[-\frac{T}{2} \log |\Gamma| - \frac{1}{2} \sum_{t=1}^T (x_t - B z_t)^T \Gamma^{-1} (x_t - B z_t)\right] \end{aligned}$$

We can rearrange the various terms by using the identity:

$$x^T A y = \text{tr}(A y x^T)$$

- Maximization respect to B :

We consider only the terms containing B , which are:

$$-\frac{1}{2} \sum_{t=1}^T \left[-\text{tr}(B^T \Gamma^{-1} x_t E[z_t^T]) - \text{tr}(\Gamma^{-1} B E[z_t] x_t^T) + \text{tr}(B^T \Gamma^{-1} B E[z_t z_t^T]) \right]$$

Maximizing respect to B we obtain:

$$\begin{aligned} \frac{\partial ELBO(q^*, \theta)}{\partial B} &= -\sum_{t=1}^T \left[-\Gamma^{-1} E[z_t] x_t^T + \Gamma^{-1} B E[z_t z_t^T] \right] = 0 \\ B &= \left(\sum_{t=1}^T x_t E[z_t]^T \right) \left(\sum_{t=1}^T E[z_t z_t^T] \right)^{-1} \end{aligned}$$

- Maximization respect to Σ :

We consider only the terms containing Σ , which are:

$$\begin{aligned}
& -\frac{T}{2} \log |\Sigma| - \frac{1}{2} E \left[z_1^T \Sigma^{-1} z_1 - \mu_0^T \Sigma^{-1} z_1 - z_1^T \Sigma^{-1} \mu_0 + \mu_0^T \Sigma^{-1} \mu_0 \right] - \\
& -\frac{1}{2} \sum_{t=2}^T E \left[z_t^T \Sigma^{-1} z_t - z_t^T \Sigma^{-1} A z_{t-1} - z_{t-1}^T A^T \Sigma^{-1} z_t + z_{t-1}^T A^T \Sigma^{-1} A z_{t-1} \right] = \\
& = -\frac{T}{2} \log |\Sigma| - \frac{1}{2} \Sigma^{-1} \left[-\mu_0^T E[z_1] - E[z_1^T] \mu_0 + \mu_0^T \mu_0 \right] - \\
& -\frac{1}{2} \sum_{t=1}^T \text{tr} \left(\Sigma^{-1} E[z_t z_t^T] \right) - \frac{1}{2} \sum_{t=2}^T \left[-\text{tr} \left(\Sigma^{-1} A E[z_{t-1} z_t^T] \right) - \right. \\
& \left. -\text{tr} \left(A^T \Sigma^{-1} E[z_t z_{t-1}^T] \right) + \text{tr} \left(A^T \Sigma^{-1} A E[z_{t-1} z_{t-1}^T] \right) \right]
\end{aligned}$$

Maximizing respect to Σ we obtain:

$$\begin{aligned}
\frac{\partial ELBO(q^*, \theta)}{\partial \Sigma} &= -\frac{T}{2} \Sigma^{-1} + \frac{1}{2} \Sigma^{-1} \left[-\mu_0 E[z_1^T] - E[z_1] \mu_0^T + \mu_0 \mu_0^T \right] \Sigma^{-1} + \\
& + \Sigma^{-1} \left(\frac{1}{2} \sum_{t=1}^T E[z_t z_t^T] \right) \Sigma^{-1} + \\
& + \Sigma^{-1} \left[\frac{1}{2} \sum_{t=2}^T \left(-E[z_t z_{t-1}^T] A^T - A E[z_{t-1} z_t^T] + A E[z_{t-1} z_{t-1}^T] A^T \right) \right] \Sigma^{-1} = 0
\end{aligned}$$

where we have used the identity:

$$\frac{\partial \text{Tr}(AX^{-1}B)}{\partial X} = -X^{-1} A^T B^T X^{-1}$$

if the matrix X is symmetric.

Now we can multiply both to left and right by Σ^{-1} , so we obtain:

$$\begin{aligned}
\Sigma &= \frac{1}{T} \left[-\mu_0 E[z_1^T] - E[z_1] \mu_0^T + \mu_0 \mu_0^T + \sum_{t=1}^T E[z_t z_t^T] + \right. \\
& \left. + \sum_{t=2}^T \left(-E[z_t z_{t-1}^T] A^T - A E[z_{t-1} z_t^T] + A E[z_{t-1} z_{t-1}^T] A^T \right) \right]
\end{aligned}$$

Also, we know that the parameter A that maximizes $ELBO[q, \theta]$ is:

$$A = \left(\sum_{t=2}^T E[z_t z_{t-1}^T] \right) \left(\sum_{t=2}^T E[z_{t-1} z_{t-1}^T] \right)^{-1}$$

and we can use the definition $\mu_0 = E[z_1]$. Plug all into the equation above, finally we obtain:

$$\Sigma = \frac{1}{T} \left[E[z_1 z_1^T] - E[z_1] E[z_1^T] + \sum_{t=2}^T \left(E[z_t z_t^T] - A E[z_{t-1} z_t^T] \right) \right]$$